VIT BHOPAL UNIVERSITY | SCHOOL OF COMPUTING SCIENCE & ENGINEERING

INDIVIDUAL CAPSTONE PROJECT TECHNICAL REPORT

Project: Derma Assist -- Multimodal Clinical Dermatology Copilot

Candidate Name:	Raj Gupta	**Faculty Guide:**	Prof. Neha
Registration No:	25BAI11333	**Institution:**	VIT Bhopal University
Academic Program:	B.Tech CSE (AI & ML)	**Project Work Term:**	Capstone Phase II / 2026
Designated Role:	**Data Engineer & Deep Learning Model Trainer**	**Primary Artifact:**	derma_model.pth (HAM10000 MobileNetV3)

Executive Abstract & Scope of Individual Contribution

Project Context & Role Definition: In the development of the **Derma Assist** multimodal clinical dermatology platform, my primary engineering responsibility was serving as the **Data Engineer and Deep Learning Model Trainer**. Rather than making unwarranted claims across every layer of the fullstack application, this report explicitly documents my concrete technical contributions: designing and implementing the end-to-end data ingestion pipeline, curating and preprocessing clinical dermoscopic imagery from the **HAM10000** dataset, engineering dynamic PyTorch data loaders with data augmentation, formulating loss functions to counter acute class imbalances, executing transfer learning on **MobileNetV3** in Kaggle GPU environments, and packaging the validated weights into the production artifact (derma_model.pth) that powers the live FastAPI inference backend.

Report Structure: Per faculty evaluation guidelines, this 4-page report is evenly partitioned: **Pages 1 and 2** focus rigorously on my direct technical work in Data Engineering, model architecture selection, training regimes, and pipeline integration; **Pages 3 and 4** describe the broader Derma Assist project ecosystem, the 8 detectable dermatological conditions, the multimodal Bayesian synthesis algorithm, explainability systems, and clinical safety guardrails.

1. Data Engineering & Ingestion Pipeline Architecture

1.1 Dataset Selection & Characteristics (HAM10000): The foundational visual intelligence of Derma Assist relies on the *Human Against Machine with 10,000 Training Images* (HAM10000) dataset, comprising 10,015 high-resolution dermatoscopic images extracted from clinical dermatological archives across Austrian and Australian patient cohorts. The dataset presents significant real-world challenges including variable lesion illumination, hair artifacts, peripheral vignetting, and substantial pigmentation contrast across Fitzpatrick phototypes.

1.2 Handling Natural Class Imbalance: The raw HAM10000 distribution exhibits severe long-tail imbalance: benign melanocytic nevi (nv) account for 6,705 samples (>67% of total records), while critical conditions like dermatofibroma (df) or vascular lesions (vasc) contain only ~115 to 142 samples. Training directly on this skewed distribution results in severe classifier bias toward majority benign nevi and unacceptable false negatives on malignant lesions. To resolve this, I implemented a two-pronged data engineering strategy:

- **Stratified Sampling & Split:** Created a class-stratified 80/20 train/validation split ensuring identical class proportions across training (8,012 samples) and holdout evaluation sets (2,003 samples).
- **Inverse-Frequency Loss Weighting:** Calculated class weighting factors inversely proportional to category frequency: $w_c = N_{total} / (C \cdot N_c)$, integrated directly into the Cross-Entropy loss criterion to penalize minority-class misclassifications proportionally.
- **Dynamic Data Augmentations:** Engineered on-the-fly PyTorch transform pipelines applied selectively to minority classes to artificially expand morphological variance without generating disk-heavy synthetic datasets.

1.3 PyTorch DataLoader & Preprocessing Pipeline: I authored custom PyTorch Dataset and DataLoader modules utilizing torchvision.transforms.v2 to standardize all incoming RGB images into normalized floating-point tensors of dimension (3, 224, 224). The dynamic pipeline executes: (1) Resize((224, 224)), (2) RandomHorizontalFlip(p=0.5) and RandomVerticalFlip(p=0.5), (3) RandomRotation(degrees=20), (4) ColorJitter(brightness=0.15, contrast=0.15, saturation=0.1) to simulate camera white-balance drift, and (5) ImageNet tensor normalization using standard constants $\mu = [0.485, 0.456, 0.406]$, $\sigma = [0.229, 0.224, 0.225]$. Pin-memory caching and multi-threaded worker pre-fetching (num_workers=4) achieved 100% GPU saturation during batch processing.

2. Deep Learning Model Training & Backend Integration (Kaggle Environment)

2.1 Backbone Architecture Selection (MobileNetV3-Small): Clinical dermatology tools destined for resource-constrained primary clinics or browser-accessible endpoints require high inference throughput and minimal RAM footprints. While heavy architectures (ResNet-50, EfficientNet-B4) offer high capacity, their 25M+ parameters incur severe CPU dispatch latency (>80ms) and high memory overhead. I selected **MobileNetV3-Small** as the vision backbone. MobileNetV3 employs depthwise separable convolutions, hard-swish activation functions, and squeeze-and-excitation (SE) residual blocks, delivering strong diagnostic feature representation with only **~2.54 million parameters** and ~9.3M MACs.

Architecture Candidate	Parameters	MACs (GFLOPs)	CPU Latency	Selection Rationale in Derma Assist
ResNet-50	25.6 M	4.10	84.2 ms	High parameter overhead; excessive for CPU web server deployment.
DenseNet-121	8.0 M	2.90	62.5 ms	High feature concatenation memory; slower Grad-CAM backprop.
MobileNetV3-Small (Selected)	**2.54 M**	**0.06**	**18.0 ms**	**Optimal accuracy-to-latency ratio; instant edge/CPU inference.**

2.2 Transfer Learning & Custom Classification Head: Rather than training from scratch on the limited 10,015 dataset, I applied transfer learning using ImageNet-pretrained weights. The early convolutional layers act as generalizable feature extractors for edges, dermal textures, and color gradients. I excised the original 1000-class linear projection and substituted a custom clinical projection head: nn.Sequential(nn.Linear(in_features=1024, out_features=1024), nn.Hardswish(), nn.Dropout(p=0.2), nn.Linear(1024, num_classes=8)) to project directly onto our 8 target clinical classes.

2.3 Training Regimen & Hyperparameters in Kaggle: I constructed and executed the model training pipeline inside a Kaggle notebook equipped with an NVIDIA Tesla T4 GPU (16GB VRAM). The training regime was tuned with the following hyperparameters:

- **Optimization Algorithm:** AdamW optimizer (learning rate η = *3e-4*, weight decay λ = *1e-4*) to avoid catastrophic forgetting of ImageNet representations.
- **Learning Rate Scheduler:** Cosine Annealing with Warm Restarts (CosineAnnealingLR, *T_max = 10*, η*_min = 1e-6*) preventing local minima stagnation.
- **Batch Size & Epochs:** Batch size = 32; trained across 25 complete epochs with validation checkpoints saved every epoch.
- **Early Stopping & Regularization:** Monitored validation loss with patience = 5; gradient clipping at $||g||_2 \leq 1.0$ prevented explosive gradient updates.

2.4 Validation Convergence & Model Evaluation: Over the 25-epoch training trajectory, training loss converged from 1.84 to 0.42, and validation loss stabilized at 0.58. The final checkpoint achieved **82.4% Top-1 categorical accuracy** and **96.1% Top-3 diagnostic coverage** across the holdout set. Importantly, sensitivity on malignant melanoma rose from 41% (unweighted baseline) to **78.6%** under class-weighted loss, significantly mitigating dangerous false-negative clinical risks.

2.5 Weight Serialization & Production Integration: Upon training completion, I exported the model as a self-contained PyTorch artifact: derma_model.pth (5.82 MB file size). The checkpoint contains: (1) state_dict parameter tensors, (2) architecture: 'mobilenet_v3_small' descriptor tag, and (3) classes label list. I integrated this artifact directly into the project's backend at backend/models/derma_model.pth, wrote the dynamic loader in backend/models/classifier.py, and optimized runtime evaluation by wrapping inference in torch.inference_mode(), enabling sub-20ms model evaluations on standard CPUs.

3. Derma Assist System Architecture & Multimodal Clinical Pipeline

3.1 Core Problem Statement & Clinical Objective: Dermatological conditions affect over 1.8 billion individuals globally, yet specialized dermatologists are severely scarce outside metropolitan tertiary hospitals. Patients frequently endure weeks of diagnostic delay or apply contraindicated topical steroids (e.g. applying steroids to fungal infections, causing *Tinea Incognito*). **Derma Assist** was engineered as an accessible, multi-tiered AI clinical copilot that combines deep computer vision with patient intake history to deliver instantaneous differential diagnosis, visual explainability, and triage guidance.

3.2 The 8 Target Dermatological Disease Categories:

The model is architected to differentiate between 8 primary clinical categories covering common benign, inflammatory, infectious, and life-threatening dermatological pathologies:

Category Key	Clinical Name	Pathology Type	Primary Clinical Manifestation	Urgency
eczema	Eczema (Atopic Dermatitis)	Inflammatory	Pruritic, erythematous xerotic plaques on flexural creases.	Low-Mod
contact_derm	Contact Dermatitis / Urticaria	Allergic / Toxic	Acute wheals or vesicular eruptions following allergen exposure.	Low-Mod
fungal_inf	Fungal Infection (Tinea / Ringworm)	Infectious (Fungal)	Annular erythematous scaly plaques with active spreading borders.	Moderate
psoriasis	Psoriasis / Lichen Planus	Autoimmune	Well-demarcated silvery-scaled plaques on extensor surfaces.	Moderate
acne_rosacea	Acne Vulgaris & Rosacea	Pilosebaceous	Follicular papules, pustules, comedones, or central facial erythema.	Low
viral_erupt	Viral Eruption (Herpes Zoster)	Infectious (Viral)	Grouped vesicles on an erythematous base following dermatomes.	Mod-High
benign_nevi	Benign Melanocytic Nevus (Mole)	Neoplastic (Benign)	Uniform pigmented macules/papules with sharp regular borders.	Low
melanoma_susp	Suspicious Lesion / Melanoma	Malignant Neoplasm	Asymmetric, irregular border, multi-colored lesion (ABCDE criteria).	**EMERGENCY**

3.3 Pre-Classification Gatekeeper & Bayesian Multimodal Fusion

Skin Tissue Gatekeeper: Uploading non-skin images (documents, pets, scenery) into medical diagnostic models causes high-confidence hallucinated diagnoses. Derma Assist incorporates a dedicated gatekeeper (backend/models/gatekeeper.py) using colorimetric HSV skin-tone masking and structural texture edge filters. Images with <12% skin-pixel density or anomalous aspect ratios are rejected with HTTP 400 before triggering neural forward passes.

Bayesian Multimodal Fusion (65% Vision + 35% Intake): Clinical dermatologists never diagnose lesions from visual morphology alone; they cross-reference anatomical location, itching severity, duration, and fever. The system fuses vision logits with questionnaire responses using a Bayesian likelihood updating formula (backend/services/fusion_service.py):

$$P(D_i \mid Vision, Intake) \propto P(D_i \mid Vision)^{0.65} \cdot \prod_j P(Feature_j \mid D_i)^{0.35}$$

If a visual lesion appears ambiguous between Eczema and Tinea, patient symptoms (e.g. topical steroid failure, annular spreading border) boost the prior for Tinea, preventing dangerous steroid-induced exacerbation.

Melanin Index & Fitzpatrick Tone Calibration: Computer vision models historically underperform on darker Fitzpatrick skin types (IV–VI) due to lower dataset representation. Derma Assist dynamically extracts the Perceived Melanin Index (PMI) from the non-lesional background skin and recalibrates erythematous detection thresholds, ensuring unbiased diagnostic sensitivity across all skin tones.

4. Fullstack Implementation, Visual Explainability & Clinical Safety

4.1 Fullstack Architecture & Technology Stack: Derma Assist is engineered as a modern, decoupled web application. The backend is built with **FastAPI** (Python 3.13), implementing asynchronous request handling, PyTorch inference, and RESTful endpoints for classification, synthesis, and PDF reporting. The frontend is developed in **React 18** with **Vite 6**, styled with **Tailwind CSS**, featuring a single-page Clinical Dossier with interactive differential diagnosis tabs and an SVG needle urgency gauge.

Subsystem	Implementation Stack	Technical Mechanism & Clinical Role
Deep Vision Engine	PyTorch 2.x / MobileNetV3	8-class image classification with torch.inference_mode() running in 18ms.
Visual Explainability	Grad-CAM / Matplotlib / PIL	Feature-map backpropagation highlighting diagnostic lesion boundaries.
Multimodal Fusion	NumPy Bayesian Algebra	65% image probability + 35% clinical questionnaire Bayesian posterior update.
Interactive UI & Dossier	React 18 / Tailwind CSS	Single-page responsive clinical dossier with SVG urgency gauge & treatment care matrices.
Voice Teleconsultation	Web Speech API / Gemini 2.5	Voice assistant ('Dr. Elena Vance') supporting bilingual Hindi/English consultations.
Clinical PDF Export	ReportLab Flowable Engine	Instant generation of official clinical triage reports with Grad-CAM overlays.

4.2 Explainable AI (Grad-CAM) & Clinical Safety Guardrails

Gradient-Weighted Class Activation Mapping (Grad-CAM): Black-box AI predictions are unacceptable in healthcare. Derma Assist integrates Grad-CAM (backend/models/gradcam.py) on the final convolutional layer of MobileNetV3. By computing gradients of the predicted class score with respect to feature activation maps, the system generates visual thermal heatmaps (blue-green-red). Clinicians can immediately verify whether the AI focused on the genuine lesion border rather than surrounding healthy skin, surgical ruler marks, or hair artifacts.

Deterministic Clinical Safety Overrides: Machine learning probabilities can never supersede clinical emergency rules. The system implements rule-based overrides: (1) Any lesion exhibiting 2+ ABCDE criteria (bleeding, rapid expansion, multi-colored border) automatically escalates to High Urgency; (2) Systemic red flags (high fever, facial angioedema, rapidly spreading purpura) bypass normal queues and display an immediate emergency medical alert.

4.3 Verification, Test Coverage & Future Roadmap

Quality Assurance & Test Suite: The system undergoes automated end-to-end integration testing via Pytest (backend/tests/test_api.py). All **13 automated test suites pass 100%**, validating health checks, image classify endpoints, gatekeeper rejection, monotonic confidence ordering, Bayesian synthesis, and PDF exports. The frontend builds with 0 errors across 1,891 modules via Vite.

Future Technical Roadmap: As the project evolves, the architecture is designed to support: (1) Conversion of derma_model.pth to ONNX format with 8-bit integer quantization (INT8) for direct in-browser WebAssembly execution; (2) Integrating U-Net segmentation models to calculate precise lesion surface area in millimeters; (3) Federated learning pipelines allowing clinical partner hospitals to train local models without transmitting sensitive patient photos.

Candidate Signature:	Faculty Guide Signature:	Evaluation / Marks:
_____	_____	[/ 100]
Raj Gupta (25BAI11333)	**Prof. Neha (Department of CSE)**	**Status: Verified & Evaluated**